

# ALF-13260 – MySQL performance issue with select\_ChildAssocOfParentByName

Status: Closed  
Project: Alfresco  
Component: Repository  
Affects Version: 4.0.d Community  
Fix Version: 4.1 Enterprise  
Security Level: External  
Environment: Debian Squeeze, Tomcat 6, MySQL 5.1.49

## 1. Problem description

In Alfresco Community 4.0.d running on Debian Squeeze with MySQL 5.1.49, heavy operations such as importing or creating thousands of documents can cause a serious performance degradation.

After such a load, the query used to resolve child associations by name, select\_ChildAssocOfParentByName, starts taking several seconds (5–6 seconds) instead of milliseconds. Since this query is executed once for each item in a folder listing, overall page response time becomes unacceptable.

Inspection of the MySQL execution plan shows that the optimizer sometimes chooses a poor plan and does not use the most appropriate index on alf\_child\_assoc, especially when statistics are out of date or skewed by recent bulk inserts.

## 2. Typical symptoms

- High CPU usage on the MySQL server during navigation or content updates.
- Slow folder browsing and very slow content creation or property updates.
- MySQL slow query log containing select\_ChildAssocOfParentByName with execution times of several seconds.
- EXPLAIN for the query showing a full scan or suboptimal index usage on alf\_child\_assoc.

## 3. Workaround: force index usage on alf\_child\_assoc

One effective workaround is to override the SQL mapping so that MySQL is explicitly instructed to use an appropriate index on alf\_child\_assoc. Example override for node-common-SqlMap.xml for the MySQL InnoDB dialect:

```
<sql id="select_ChildAssoc_FromSimple_By_Parent_Index"> from alf_child_assoc assoc use index
(idx_alf_cass_pri) join alf_node parentNode on (parentNode.id = assoc.parent_node_id) join alf_store
parentStore on (parentStore.id = parentNode.store_id) join alf_node childNode on (childNode.id =
assoc.child_node_id) join alf_store childStore on (childStore.id = childNode.store_id) </sql> <select
id="select_ChildAssocOfParentByName" parameterType="ChildAssoc" resultMap="result_ChildAssoc">
<include refid="alfresco.node.select_ChildAssoc_Results"/> <include
refid="alfresco.node.select_ChildAssoc_FromSimple_By_Parent_Index"/> where parentNode.id =
#{parentNode.id} and assoc.child_node_name = #{childNodeName} and assoc.child_node_name_crc =
#{childNodeNameCrc} </select>
```

This change forces MySQL to use the specified index for the association lookup. Several installations reported stable performance for weeks after applying this patch.

## 4. Alternative: create a dedicated index and hint MySQL

In some large repositories, creating an additional index on child\_node\_name and child\_node\_name\_crc and then forcing MySQL to use that index provides better performance. Example index creation on alf\_child\_assoc:

```
CREATE INDEX idx_alf_cass_pri2 ON alf_child_assoc (child_node_name, child_node_name_crc);
```

The SQL mapping can then be adjusted to use the new index:

```
<sql id="select_ChildAssoc_FromSimple_By_Parent_Index"> from alf_child_assoc assoc use index (idx_alf_cass_pri2) join alf_node parentNode on (parentNode.id = assoc.parent_node_id) join alf_store parentStore on (parentStore.id = parentNode.store_id) join alf_node childNode on (childNode.id = assoc.child_node_id) join alf_store childStore on (childStore.id = childNode.store_id) </sql>
```

Reports from production environments with more than a million documents indicate that this approach reduces per query execution time from around 5 seconds to well under a second, restoring acceptable performance for content creation and updates.

## 5. Core fix: include type\_qname\_id to leverage composite index

There is a composite index on alf\_child\_assoc that includes parent\_node\_id, type\_qname\_id, child\_node\_name\_crc, and child\_node\_name. However, the original query did not restrict type\_qname\_id, which prevented MySQL from fully using this index. The effective long term fix is to include type\_qname\_id in the WHERE clause. Conceptually, the change is: Original pattern:

```
<select id="select_ChildAssocOfParentByName" parameterType="ChildAssoc"
resultMap="result_ChildAssoc"> <include refid="alfresco.node.select_ChildAssoc_Results"/> <include
refid="alfresco.node.select_ChildAssoc_FromSimple"/> where parentNode.id = #{parentNode.id} and
assoc.child_node_name = #{childNodeName} and assoc.child_node_name_crc = #{childNodeNameCrc}
</select>
```

Updated pattern:

```
<select id="select_ChildAssocOfParentByName" parameterType="ChildAssoc"
resultMap="result_ChildAssoc"> <include refid="alfresco.node.select_ChildAssoc_Results"/> <include
refid="alfresco.node.select_ChildAssoc_FromSimple"/> where parentNode.id = #{parentNode.id} and
assoc.type_qname_id = #{typeQNameId} and assoc.child_node_name = #{childNodeName} and
assoc.child_node_name_crc = #{childNodeNameCrc} </select>
```

By constraining type\_qname\_id, the optimizer can make use of the existing composite index, which dramatically improves performance without relying on explicit index hints. Internal testing for large scale site creation scenarios showed a reduction from tens of seconds per site creation to approximately 150 ms.

## 6. Validation and recommended approach

After applying the fix that includes type\_qname\_id in the query, performance tests on later builds (4.0.2(3) and 4.1.0) no longer showed slow queries for select\_ChildAssocOfParentByName and overall system performance under heavy load was acceptable.

Recommended approach:

1. If you need an immediate workaround on an affected version, consider:
  - Adding an explicit index hint as shown in section 3, or
  - Creating a dedicated index and forcing its usage as in section 4.
2. For a long term solution, ensure that your Alfresco version includes the change that adds assoc.type\_qname\_id to the WHERE clause of select\_ChildAssocOfParentByName, or apply an equivalent patch.

3. After any change, run EXPLAIN on the query and monitor the MySQL slow query log under representative load to confirm that the appropriate index is being used and that execution time is acceptable.