




## a. Overview



Data Modeling




Cost Reduction




Performance



Time to live



Caching



Local Development Environment

## b. Relational vs. NoSQL Data Modeling

Normalized (3NF)

No data duplication

Less disk usage

More CPU usage

Schema and queries are separated

Allows flexible queries

Lower performance

and non-scalable

Data can be denormalized

May have duplication

more disk space


lower CPU usage

Schema design based on queries


Not flexible

Higher performance

# RDBMS



No SQL → Does not stop from using multiple tables



Single table design is possible

## Trade Offs of NoSQL Data Modeling

Get all relate data with  
one request  
low latency

understand access patterns  
Requires thinking about queries upfront  
invest time in design

## C. 1:1 and 1:M relationships

### Modeling

#### 1:1 relationships

2 tables



Advantages (or) Why do this?

- DynamoDB item size limitation  $< 400$  KB
- Updating smaller item costs less. (even for updating one attribute)
- Can create more indexes


#### 1:M relationships

Partition  
key  
eg. Order  
Id


Sort key  
eg. ItemId

eg.  
We can get  
all items in  
an order.

# Representing Hierarchical Data



A different sort key structure

CITY # ZONE # DEPT.

COUNTRY	DEPARTMENT	EMPLOYEE COUNT
India	Chennai # South # Finance	25
India	Chennai # North # Sales	33
India	Mumbai # South # HR	5

Partition Key

Sort Key

Attributes


Query → Get all departments

COUNTRY = India

DEPARTMENT Begins with  
"Chennai"

## d. Representing Many-to-many relationships

(eg) Students Courses



RDBMS


STUDENTS


COURSES


STUDENTSTO  
COURSES

Association  
table

Query uses joins

```
SELECT * FROM STUDENTS AS S
```

```
JOIN STUDENTSTOCOURSES STC
```

```
ON S.Id = STC.StudentId
```

```
JOIN COURSES C ON C.Id = STC.CourseId
```

For M:M in DynamoDB (NoSQL)  
we just need a single table



# NoSQL (DynamoDB)


Other attributes

	Id	sk	Attributes
Student {	s1	s1	Name: Srihari
Course {	c1	c1	Name: DynamoDB ...
Student to Course {	s1	c1	Grade: A ...

Similarly we can have data for other students and courses

Note:

DynamoDB is schemaless



Important → Use Global Secondary Index

to model many to many

Normal table

It was from the perspective of students

GSI

It is from the perspective of courses

NOTE: DynamoDB creates copy of data in a

new table for a GSI. SRIHARI SRIDHARAN

Get all info about students } Single query  
Get all info about courses }

No Joins → Efficient

Key take aways

- One table instead of many.
- Different types in one table.
- Single attribute used for multiple fields.

e. Concatenated attributes

{ use case → query all orders  
delivered in July 2021  
→ concatenated attributes

OrderId	UserId	Status	Date	Status_Date
1	1 ↑ GSI PK	DELIVERED SHIPPED CANCELLED DELIVERED	2021-01-02 2021-05-03 2021-07-01 2021-07-01	This column stores the status and date as concatenated string

Query: Status\_Date BEGINS-WITH  
"DELIVERED#2021-07"

## f. Hot keys

### What ?

Some partitions  
getting uneven  
load - R/W

### How to deal with them?

- Choice of partition keys
- Adaptive Capacity
- Shard Limits

## Avoid hot keys

- Ensure uniform load on all partitions

↳ Hard to achieve

- Use caching

- Select good partition key

### How to select a partition key ?

x Boolean ? Bad choice (only 2 values)

x Numeric values → limited range → Bad.

✓ UUID - good choice for a partition key

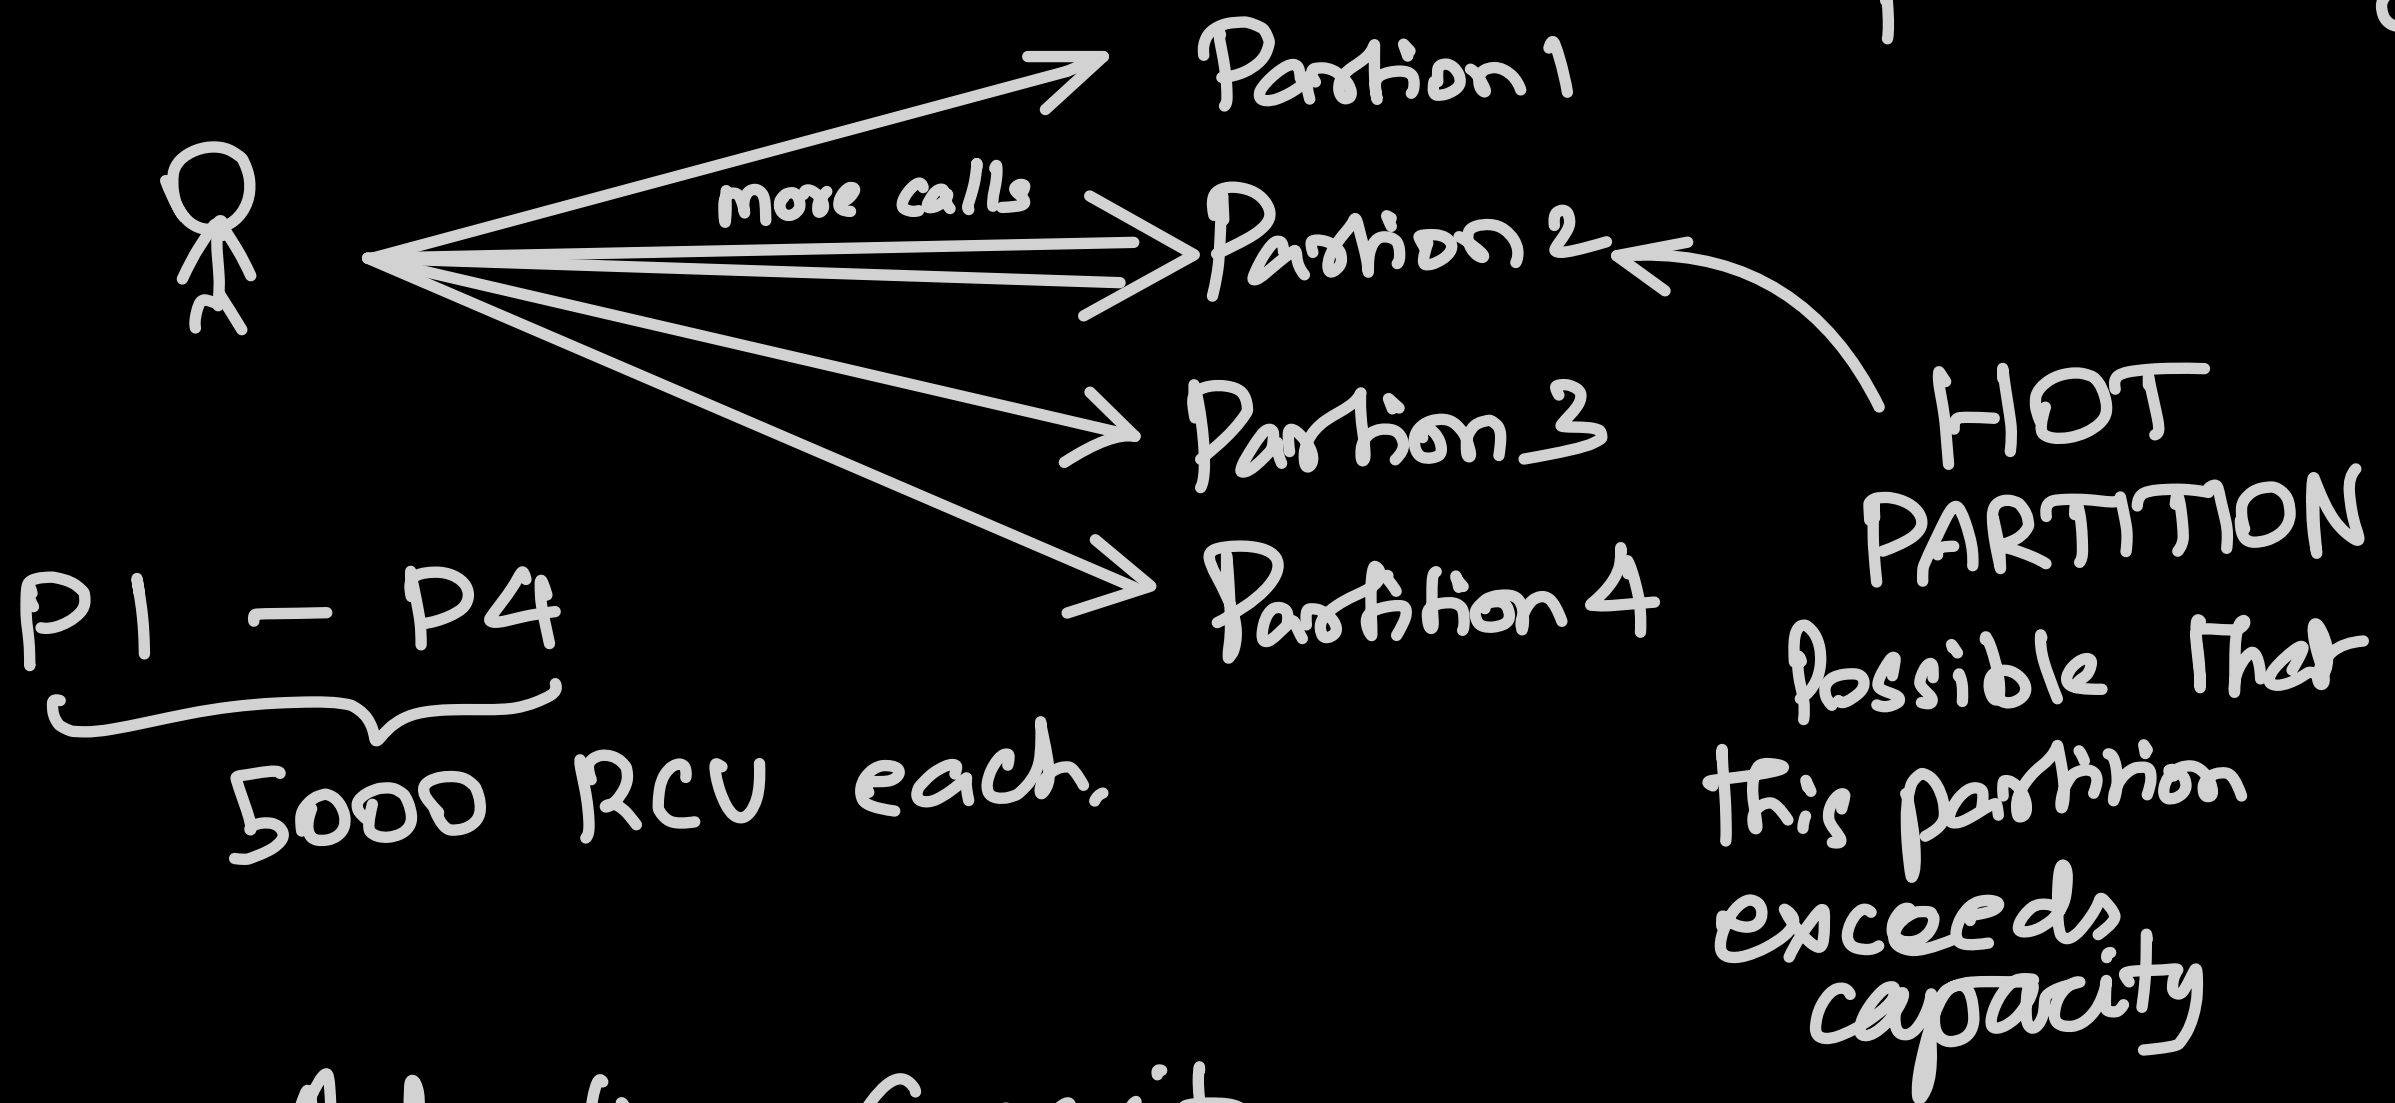
OK ✓ Numbers - You need a good range of




A table with 20000 RCU - assume it receives read requests of 5000 RCU

Anything to worry?

Total 20000 RCU → If depends on partitioning



Adaptive Capacity



It is hard to achieve even load and DynamoDB shifts some capacity to partitions under load.

Works if consumed capacity is less than the provisioned capacity.

## Shard Limits

Even with adaptive capacity there are total limits. It will work as long as total consumed RCUs is less than 3000 RCUs and WCUs is less than 1000 WCUs on a single partition.

## g. Write Sharding

Some items can be popular than the others → Can receive too many write requests

eg: Election Voting

↳ Some candidates are more popular - may need more than 1000 WCUs.

Solution: Split into multiple items and aggregate in the application.

Random Suffixes

Calculated Suffixes

## Random Suffixes

Key = Id + rand(10)


DynamoDB

eg

candidate 42 \_ 1 →

candidate 42 \_ 7 →

candidate 42 \_ 10 →



Prefix

Random  
Suffix

(Application

reads the data  
and aggregates it)

## Calculated Suffixes

key = date() + order-Id() + rand(10)

→ eg: 21.07.2021 # 23 # 9

21.07.2021 # 23 # 7

## h. Reducing costs

As costly as we use RCUs and WCU's

★ Send less data to DynamoDB

- Store big items in S3
- Use data compression
- Attributes Projection
- Split big items into multiple tables

★ Exploit temporal access patterns

for time series data

◦ Recent Data → More RCUs and WCVs

◦ Older Data → Lesser RCUs and WCVs

◦ Very old data → Move to S3



i Using Scans appropriately

Avoid scans → Use queries instead

Avoid spikes in read requests

from scans → Can affect production traffic

Read using small page sizes

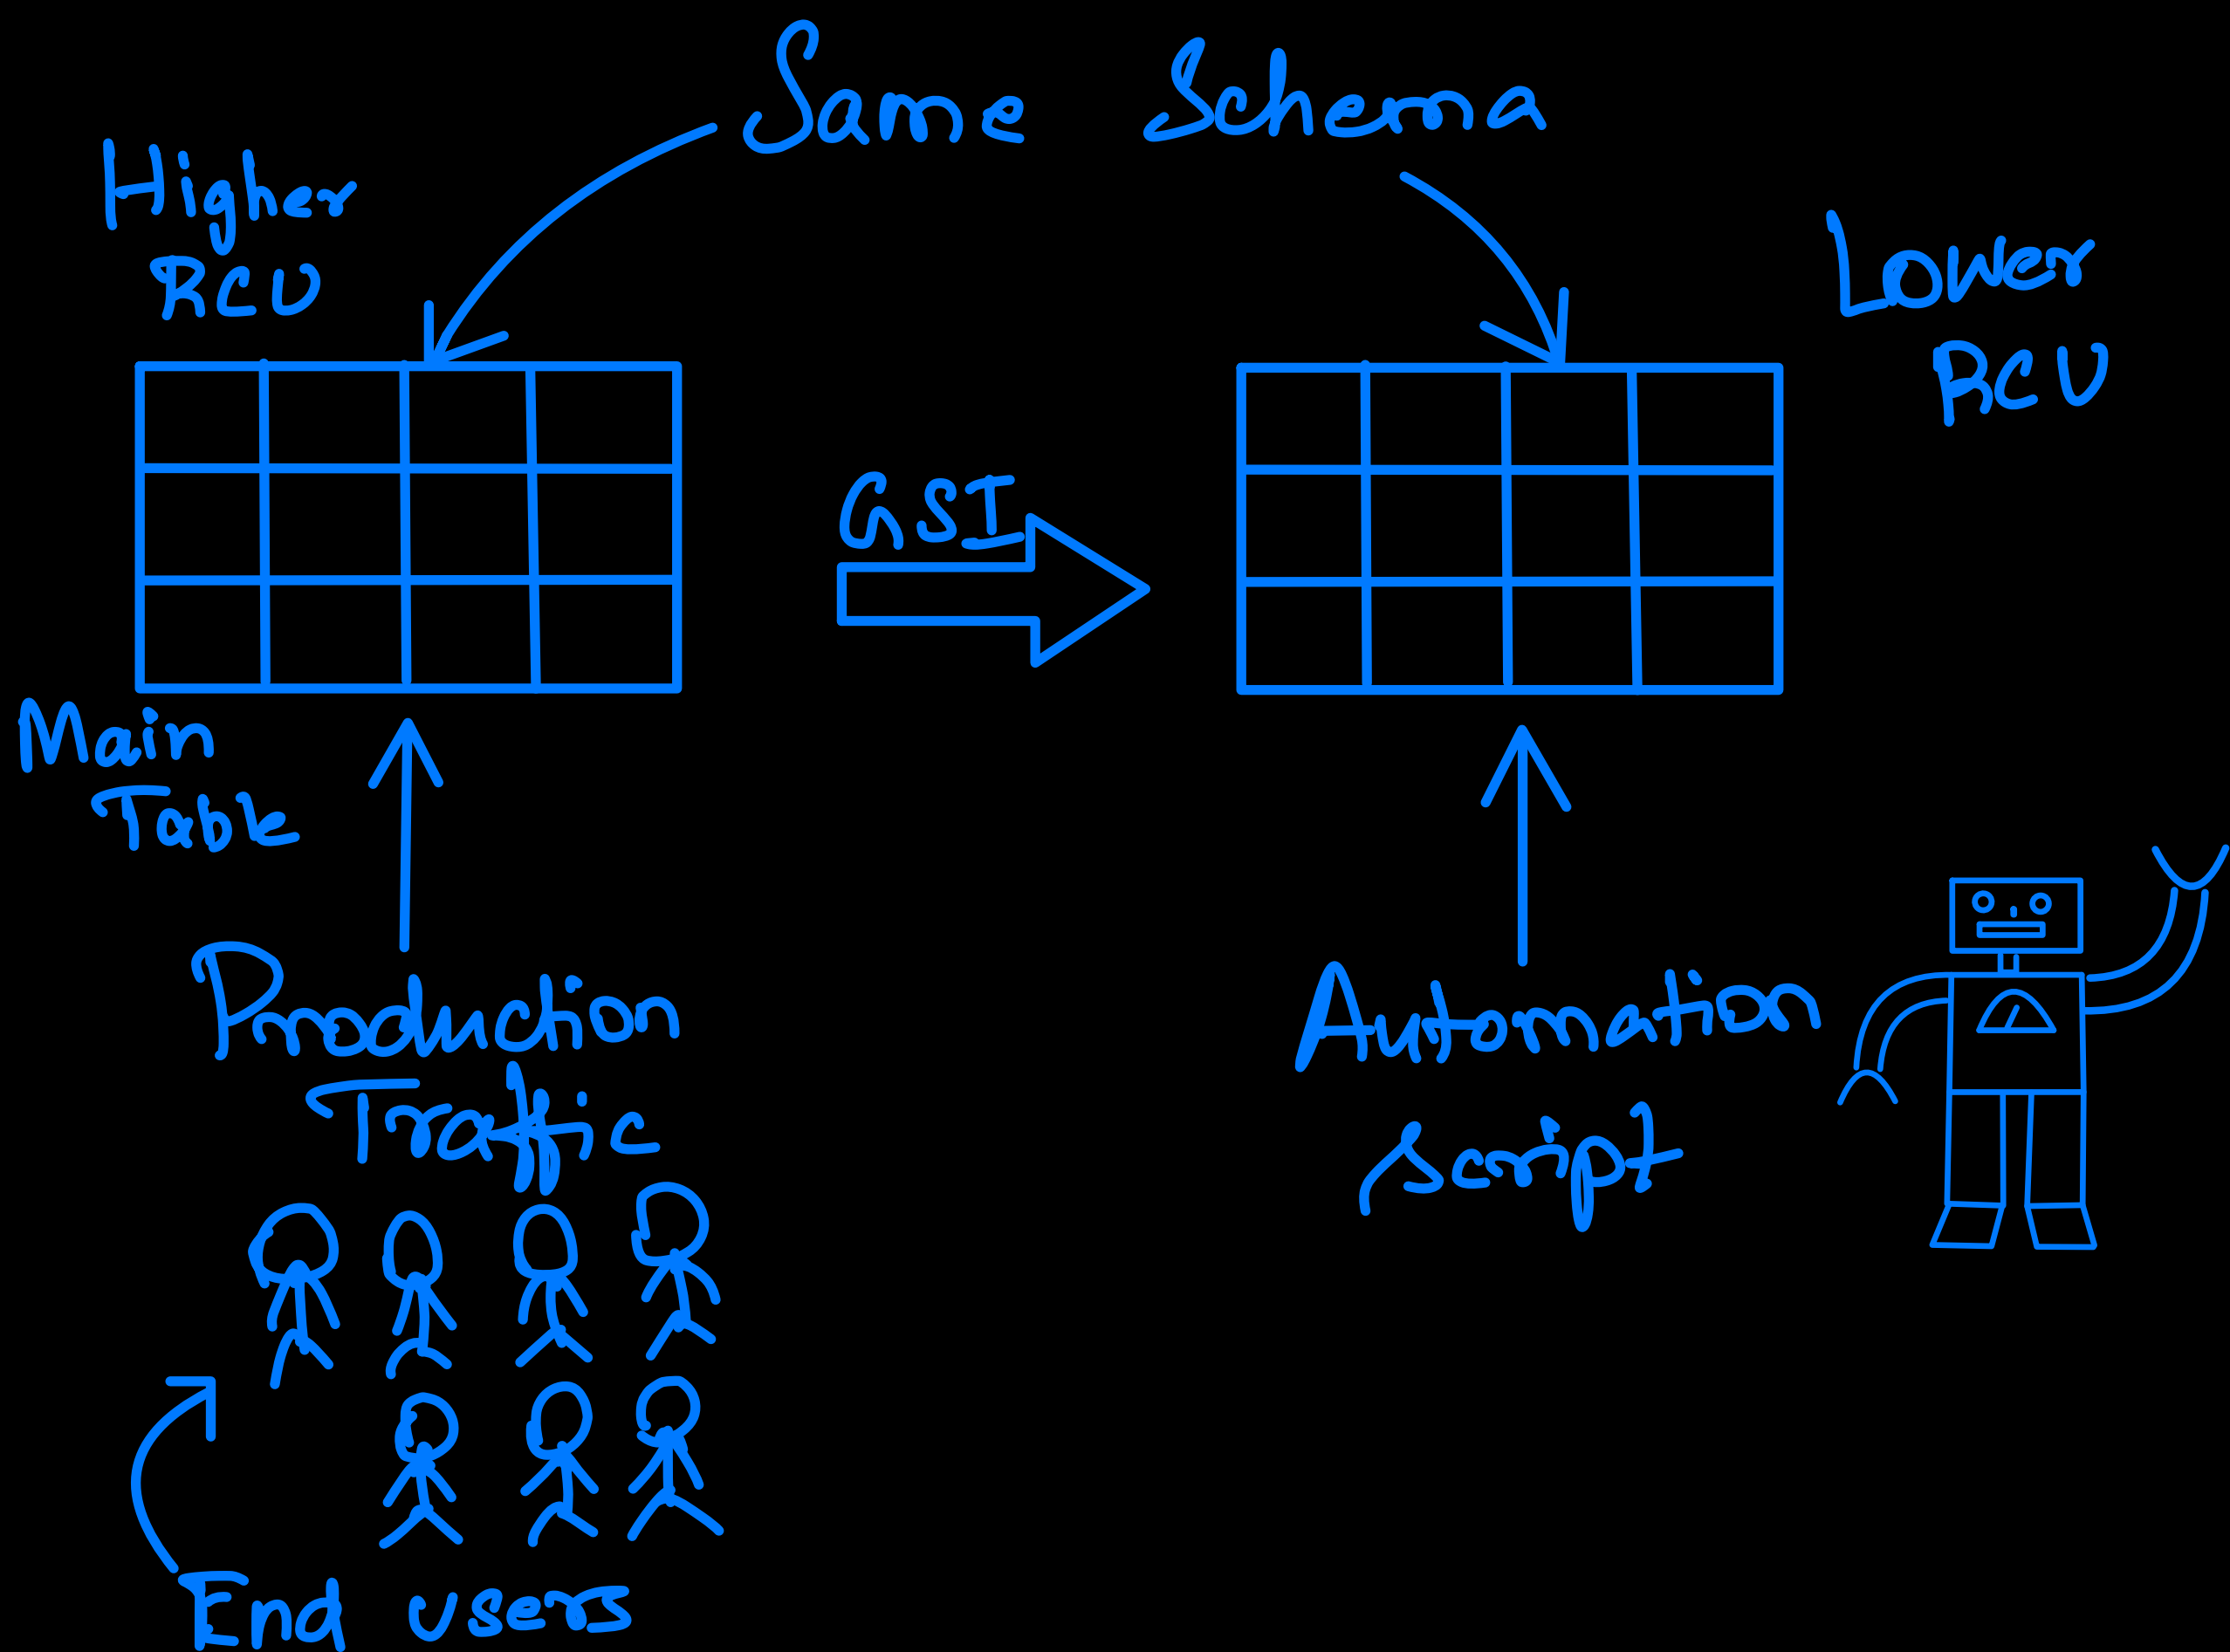
Pause between reading chunks of data

Parallel Scans → Allows to read data faster

↳ Does not put all load on a single partition.

SRIHARI SRIDHARAN

# Shadow Table



## j Time-to-live

DynamoDB removes item from the table when they expire

No extra cost

Allows to save money

No need to manually remove items



## TTL use cases

- Legal requirements - Don't use TTL when you need to retain data for legal and compliance
- Session Data
- Temporary Data

TTL can be set at the table level


## K. Sparse Indexes

Filter a small subset of items

↳ Have a specific criteria

Have a GSI

on that attribute



Only items with this attribute are copied

Not all the items in the main table have this attribute.

## l. Caching with DynamoDB

Why? Avoid hot partitions  
Decrease latency

Read - intensive applications


SRIHARI SRIDHARAN

# DAX - DynamoDB Accelerator



- In-memory acceleration
- sub millisecond latency
- Fully managed
- Scalable up to 10 nodes
- Sits between DynamoDB and the application
- Easy to use and same API

## Before DAX

Application



Cache




DynamoDB

Application managed Caching 😞

This is how DAX Works

After an item is saved in DynamoDB it gets cached in DAX.

Application



If item is in itemcache it is returned, if not it is obtained from DB, stored in item cache and returned to the app.

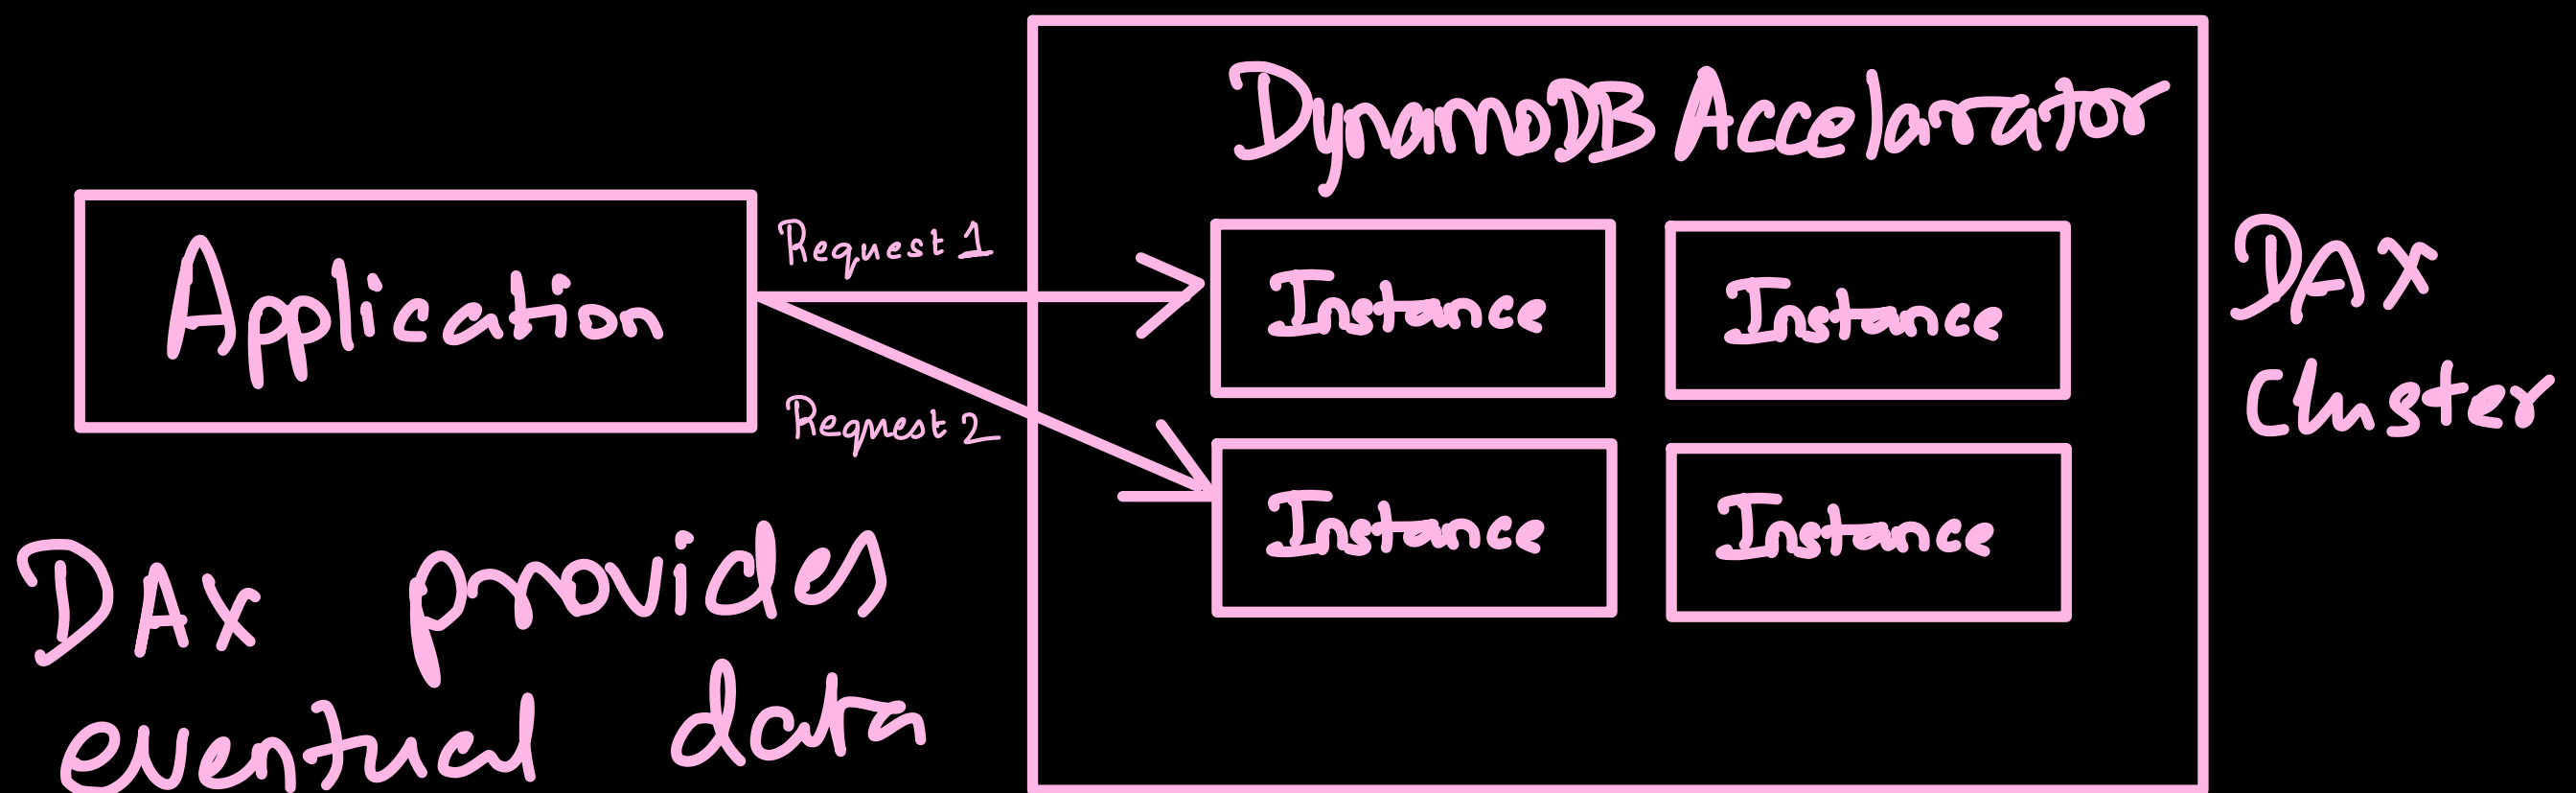
DynamoDB Accelerator (DAX)

Item Cache

Query Cache



DAX has multiple instances



DAX provides eventual data

consistency "If your application requires strong consistency you cannot use DAX"

NOTE: Application can bypass DAX and bulk write into DynamoDB, just that items will not be cached.  
A DAX cluster can only be accessed by an application that is running in the VPC.

m. DynamoDB and AWS Lambda

λ Lambda - FaaS, serverless,


↳ can process DynamoDB streams

Each shard can be processed by an instance of λ

## n. Local Development Environment

Why run locally?

- Fast
- Convenient
- Low cost
- Tests can be run at the same time



2 options - DynamoDB Local , Docker Image

Thanks for  
reading !!

THE  
END

All the  
best !!

Warm regards  
SRIHARI SRIDHARAN