



JavaOneSM
Sun's 2000 Worldwide Java Developer Conference™

Performance Tuning for the JDBCTM API

What Works, What Doesn't, and Why.

Mark Chamness
Sr. Java Engineer
Cacheware

Overall Presentation Goal

Illustrate techniques for optimizing JDBC API-based calls from the Java platform, including examples for an Oracle database implementation.

Learning Objectives

As a result of this presentation, you will:

- Design better JDBC implementations
- Recognize potential performance bottlenecks
- Make more money

Speaker's Qualifications

- Professional Java developer since 1995, including 3 years at Sun Microsystems.
- Implemented JDBC for production systems using the major databases: Sybase, Informix, and Oracle.
- Worked on more than 3,000 Java bugs at Sun Microsystems.



Agenda

- 1) Why optimize?
- 2) Basic API techniques.
- 3) Design Strategies.
- 4) Advanced Driver Tuning methods.

Why Optimize?

- On average, a web request performs 4 database queries.
- Experience has shown that database calls are typical performance bottleneck.
- Bad JDBC can overwhelm the database.

JDBC API

Most Versatile



- SQL: "SELECT * FROM TABLE"
- `java.sql.PreparedStatement`
- `java.sql.CallableStatement`
- Cache data on client.

Most Optimized

JDBC API

SQL Statements

- Most flexible
- Least reliable
- Must be recompiled in database for each use

JDBC API

PreparedStatement

- Represents a precompiled SQL statement
- Can be used to efficiently execute statement multiple times
- Somewhat flexible – can create new ones as needed

JDBC API

```
PreparedStatement pstmt =  
    con.prepareStatement("UPDATE EMPLOYEES  
    SET SALARY = ? WHERE ID = ?");
```

```
pstmt.setBigDecimal(1, 153833.00);
```

```
pstmt.setInt(2, 110592);
```

```
pstmt.execute();
```

JDBC API

java.sql.CallableStatement

- Used to execute SQL stored procedures.
- Same syntax as PreparedStatement.
- Least flexible.
- Most optimized DB call.

JDBC API

Cache

- Keep data within client to reduce the number of round-trips to the database.
- Lesson: The less JDBC the better.

Basic Design Techniques

Use Database Connection Pool

- Don't use `DriverManager.getConnection()` often. JDBC connections can take 0.5 to 2 seconds to create.
- Create Pool of Connections and reuse them.
- Necessity for any production system.

Basic Design Techniques

Use multi-threading with Connection Pooling to address network latency:

- Threads can issue queries over separate database connections.
- This improves performance to a point.

Basic Design Techniques

Single-batch Transactions

Collect set of operations and submit transaction in one statement:

```
BEGIN TRANSACTION
  UPDATE TABLE1...
  INSERT INTO TABLE2...
  DELETE TABLE3
COMMIT
```

Basic Design Techniques

Single-batch Transactions

- DB obtains necessary locks on rows and tables, uses and releases them in one step
- Depending on transaction type, separate statements and commits can result in more DB calls and hold DB locks longer



Basic Design Techniques

Single-batch Transaction Types

Significantly different effects!

`java.sql.Connection`

- TRANSACTION_READ_COMMITTED
- TRANSACTION_READ_UNCOMMITTED
- TRANSACTION_REPEATABLE_READ
- TRANSACTION_SERIALIZABLE

Basic Design Techniques

Don't have transaction span user input

- Application sends BEGIN TRAN and SQL, locking rows or tables for update
- Application waits for user to press key before committing transaction

Basic Design Techniques

Solution: Optimistic locking

- Optimistic locking employs timestamps and triggers in queries and updates
- Queries select data with timestamp values
- Prepare a transaction based on that data, without locking data in a transaction



Basic Design Techniques

Smart Queries

- Make queries as specific as possible
- Put more logic into SQL statements
- DB are designed to use SQL efficiently
- Proper use of SQL can avoid performance problems

Basic Design Techniques

Smart Query Ex: get employees in ENG dept

Instead of:

```
SELECT * FROM employees;
```

```
SELECT * FROM dept;
```

(and joining on Java application side)

Use database join:

```
SELECT employees.* FROM employees E,  
dept D WHERE E.DEPTNO = D.DEPTNO AND  
D.DEPTTYPE = 'ENG';
```



Basic Design Techniques

Smart Queries

- Minimize ResultSet before crossing network
- Many performance problems come from moving raw data around needlessly

Basic Design Techniques

Smart Query Guidelines

- Use DB for filtering
- Use Java for business logic
- DB does filtering very well
- DB business logic is poor
- (At least very inconsistent between database vendors.)

Basic Design Techniques

Keep operational data set small as possible

- Move non-current data to other tables and do joins for rarer historical queries
- Otherwise, index and cluster so frequently used data is logically and physically localized

Advanced Driver Tuning

- Special options for each JDBC driver
- No common standard
- Improve performance by reducing round trips to the database.
- Ex. Oracle driver performance extensions

Advanced Driver Tuning

Oracle Performance Extensions

- 1) Row Prefetch
- 2) Batch Updates
- 3) Suppress DatabaseMetaData
TABLE_REMARKS Columns
- 4) Specify Column Types

Advanced Driver Tuning

1. Row Prefetch

- Use client-side buffers
- Replace round trips by local manipulation of rows returned by query
- Use `OracleStatement.setRowPrefetch()`

Advanced Driver Tuning

2. Batch Updates

- Reverse Prefetch
- Does for data going to DB what prefetch does for data coming from it
- `OraclePreparedStatement.setExecuteBatch`

Advanced Driver Tuning

2. Batch Updates

- Standard JDBC makes a trip to DB for each `PreparedStatement.executeUpdate`
- Batching: When number of queued requests reaches batch size, sends them to DB

Advanced Driver Tuning

Suppress DatabaseMetaData
TABLE_REMARKS Columns

- Avoids expensive outer join operation
- DatabaseMetaData.getTables() & getColumnns() are slow
- OracleConnection.setRemarksReporting()

Advanced Driver Tuning

4. Specify Column Types

- Standard JDBC: 1st trip to DB used to determine column types of ResultSet
- Then converts data to requested return type (if necessary)

Advanced Driver Tuning

4. Specify Column Types

- Specify column types, JDBC makes one fewer trip to DB
- The server performs any necessary type conversions (which is optimized)

Advanced Driver Tuning

4. Specify Column Types

Must specify data type for each column of expected ResultSet

1. `OracleStatement.clearDefines()`
2. `OracleStatement.defineColumnType(..)`
3. `OracleStatement.executeQuery()`

Summary

Optimization Stages

- 1) Leverage the strengths of the DB
- 2) Use the full range of java.sql API
- 3) Design for Performance – Connection Pools, Multi-Threading, etc.
- 4) Implement driver performance extensions



JavaOneSM
Sun's 2000 Worldwide Java Developer Conference™

Q&A